

Welcoming the Era of Deep Neuroevolution

Original article by Uber Data:

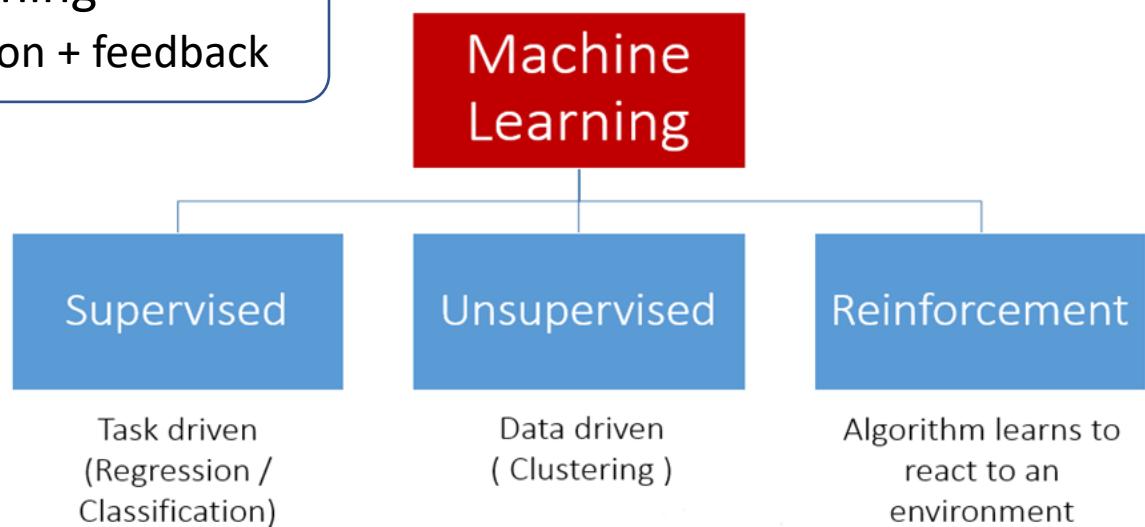
<https://eng.uber.com/deep-neuroevolution/>

Presented paper – Genetic algorithms for deep learning:

<https://arxiv.org/abs/1712.06567>

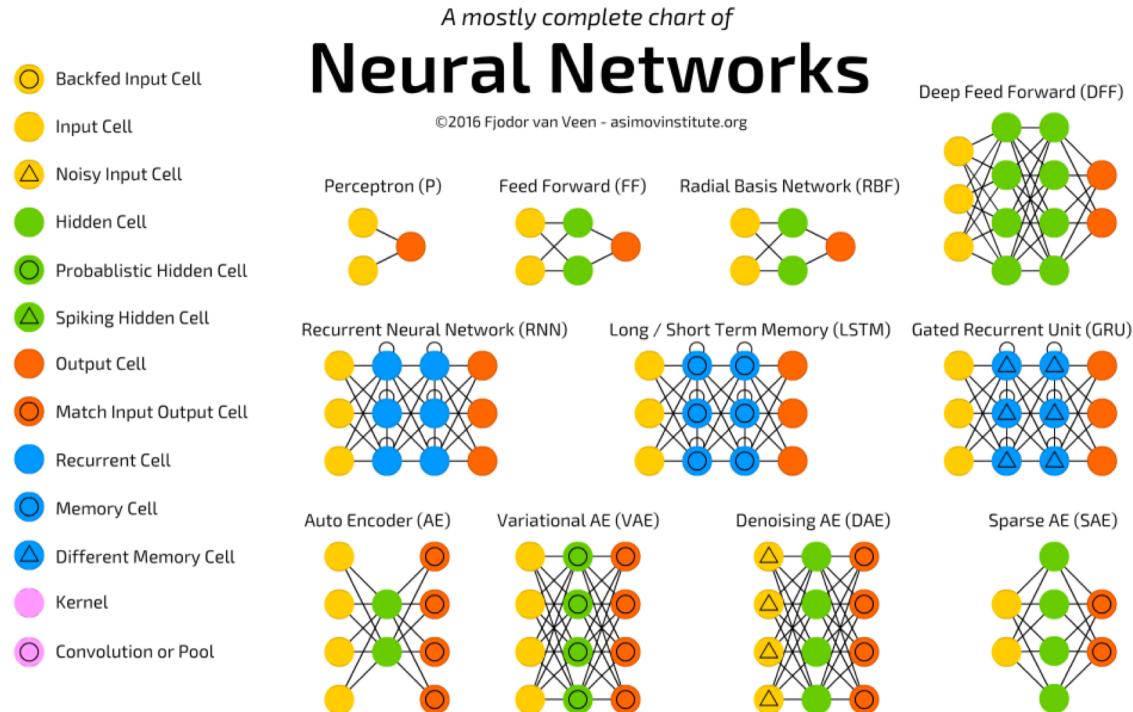
Types of machine learning

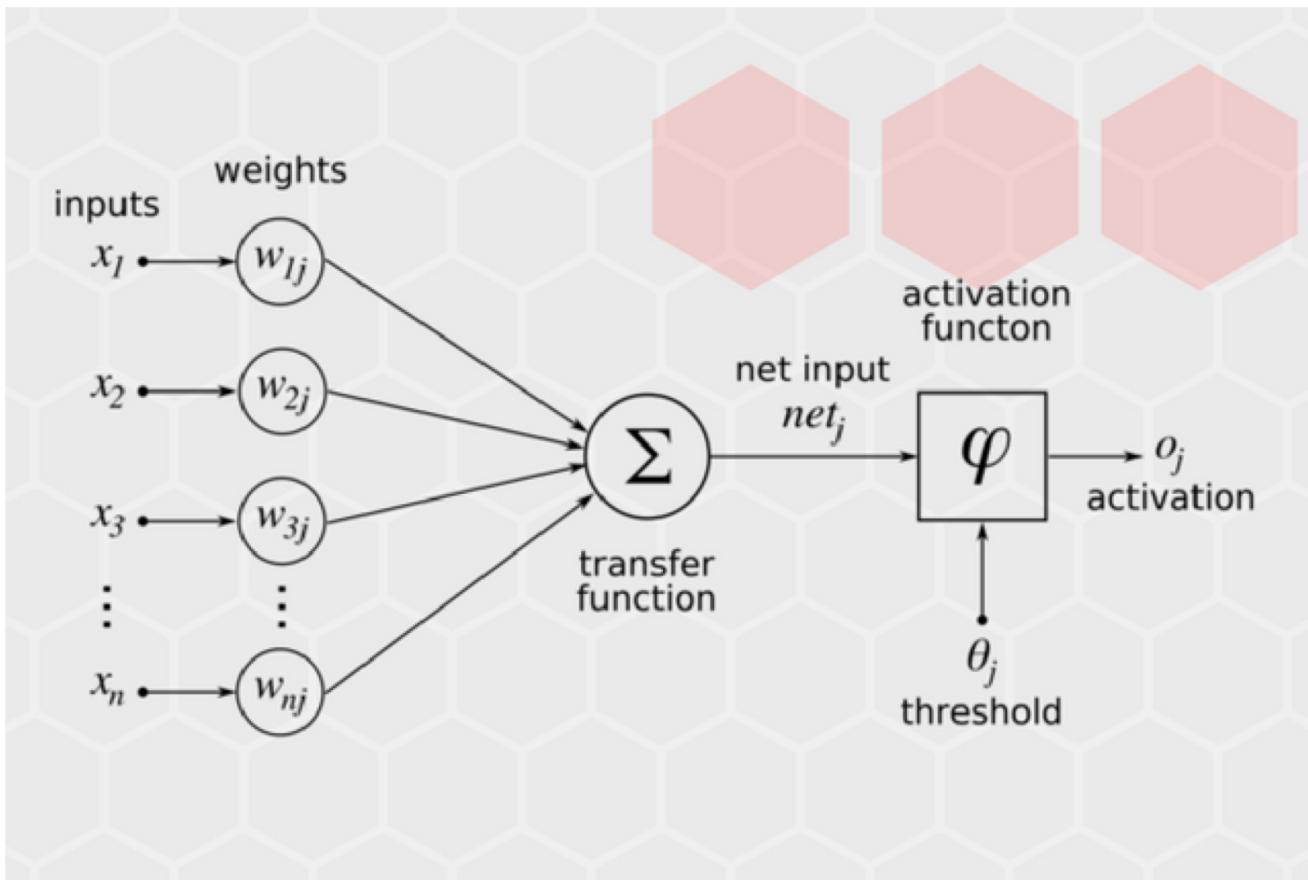
- **Supervised** learning
 - we know (input, output) data
- **Unsupervised** learning
 - we only know the (input) data
- **Reinforcement** learning
 - we know observation + feedback

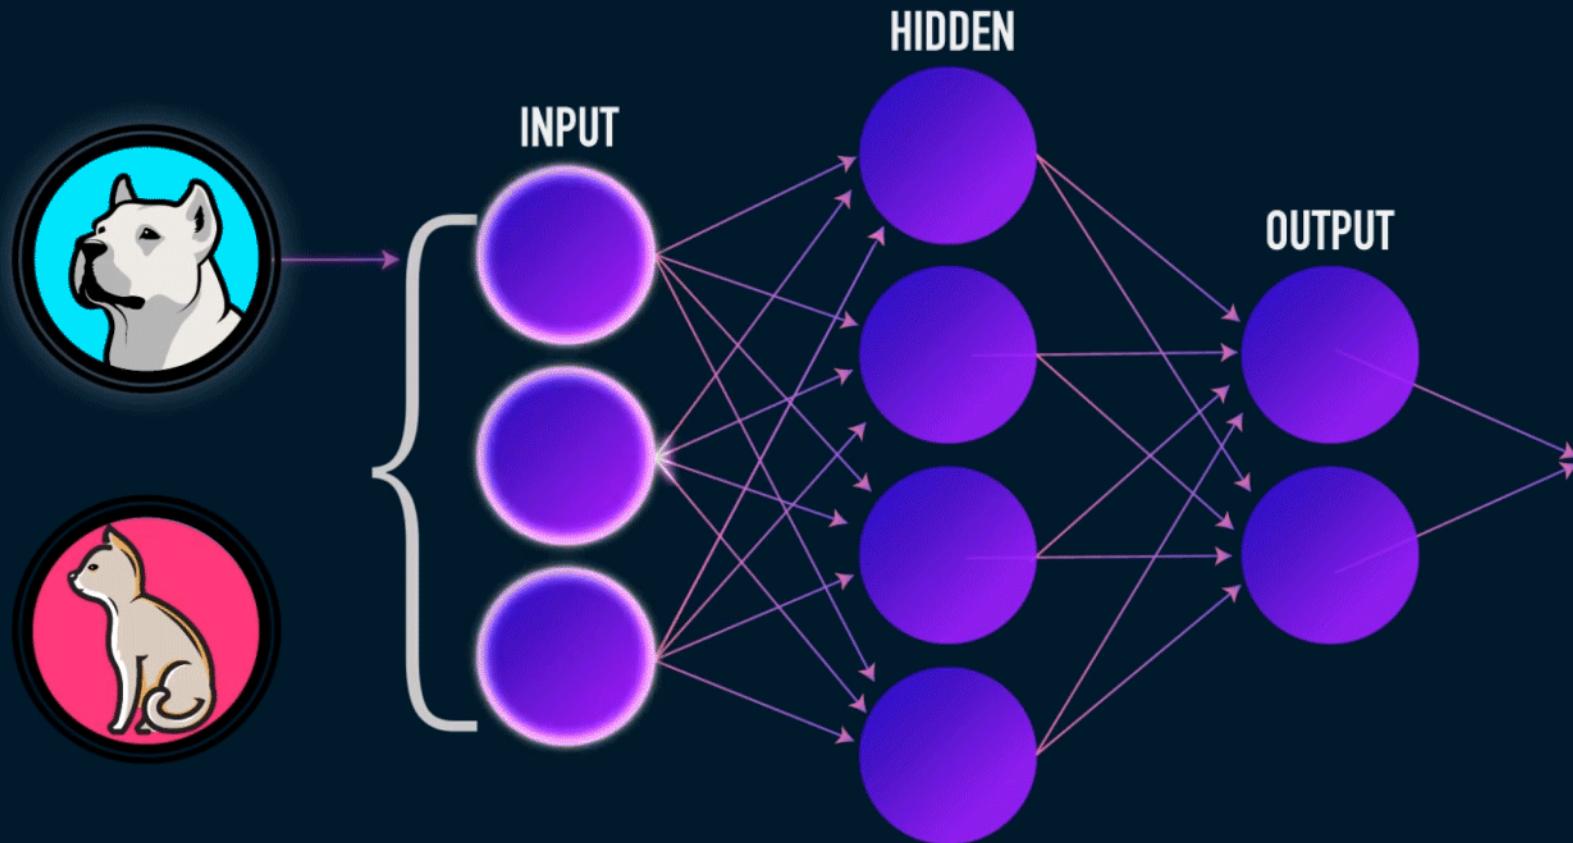


But what is a Neural Network (NN)?

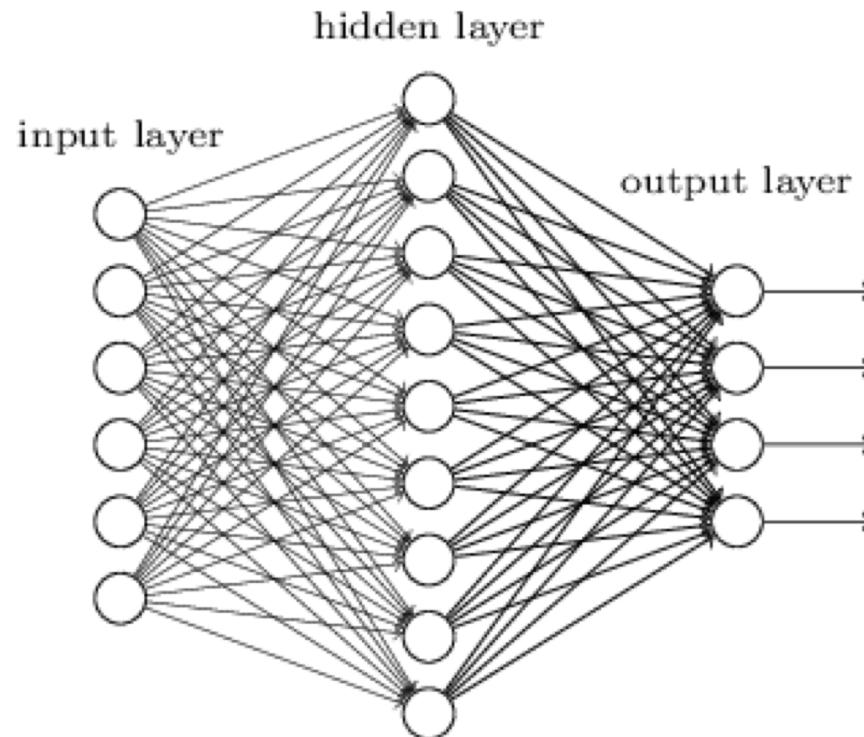
- Universal approximator. Inspired by human brain.
- It can do awesome things!







Is this a deep Neural Network?



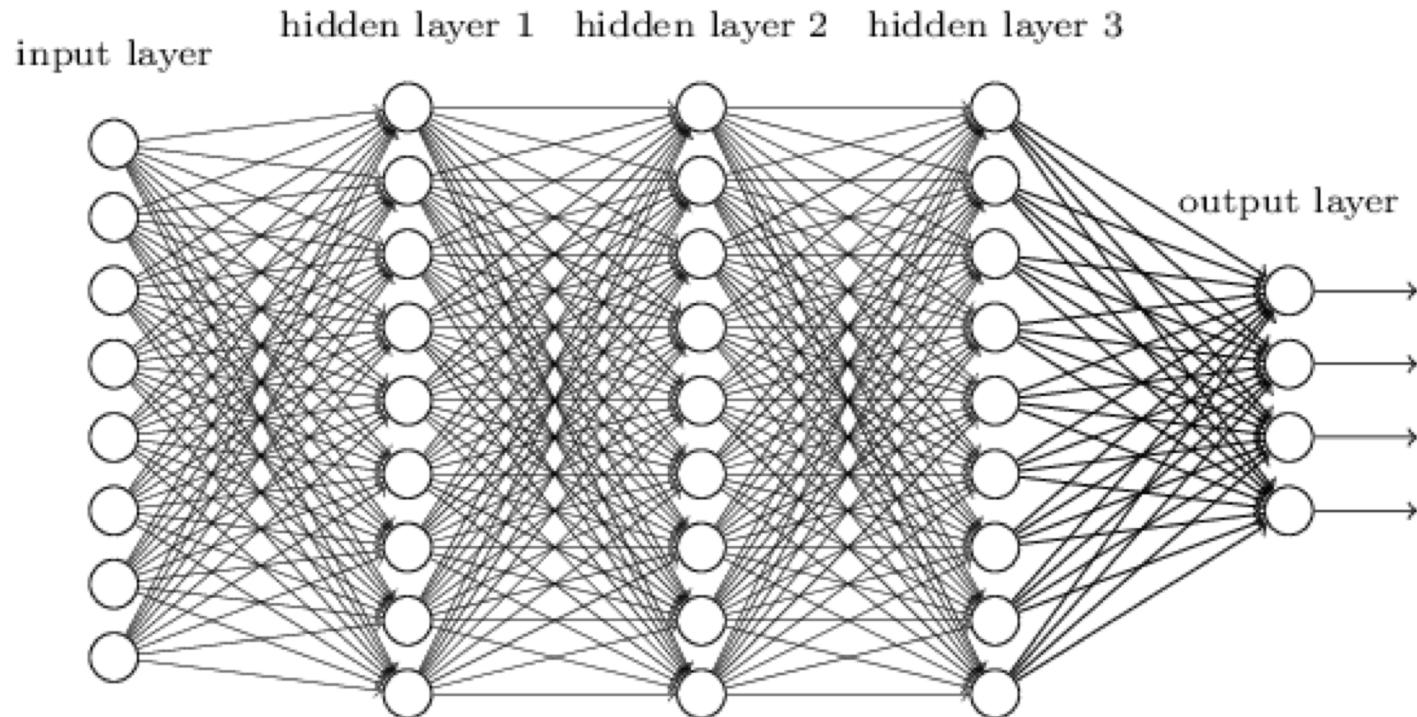
well, it does not look very deep

A scene from the movie "The Wolf of Wall Street". Leonardo DiCaprio's character, Jordan Belfort, is seated at a table in a restaurant, looking down at his food. He is wearing a dark suit and tie. Jonah Hill's character, Donnie Azoff, is seated across from him, partially visible on the right side of the frame. The lighting is warm and dramatic. The text "WE NEED TO GO DEEPER" is overlaid in large, bold, white letters.

WE NEED TO GO

DEEPER

Is this a deep Neural Network?



.. nobody really knows

How do Neural Networks learn?

Neural network is basically a function:

$$NN(x^1, x^2, x^3 \dots x^n) = (y^1, y^2, y^3 \dots y^m)$$

The inputs and outputs are given, so what do we adjust?
weights!

How do we know how much should the weights to be adjusted?
We calculate how „wrong“ the NN is.

This is usually called a loss function: on all data we sum the error.

$$e.g. MSE = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

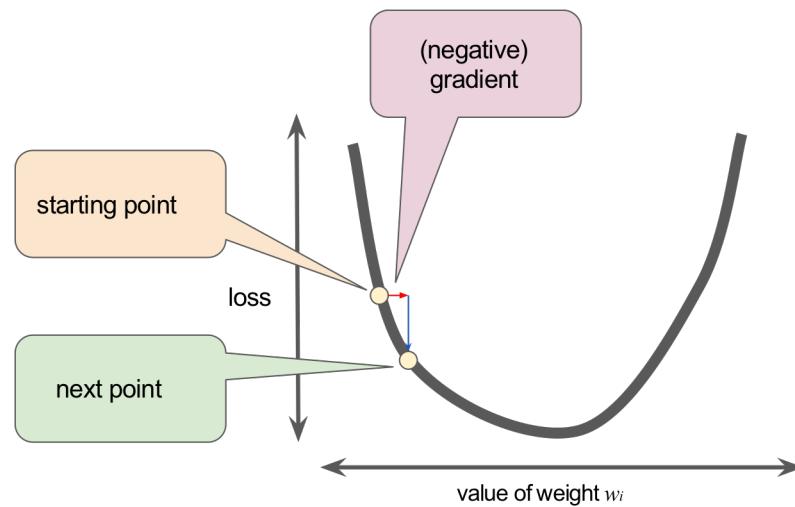
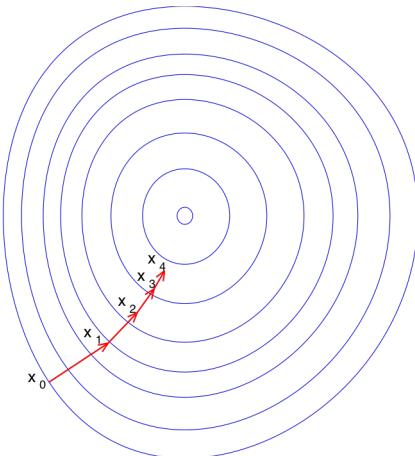
And we try to **minimise** this loss function **by adjusting weights**.

How do we minimise the loss function?

- we get the **gradient!** – “vector” that shows in which way we adjust the weights to minimize a loss function

Gradient Descend (GD)

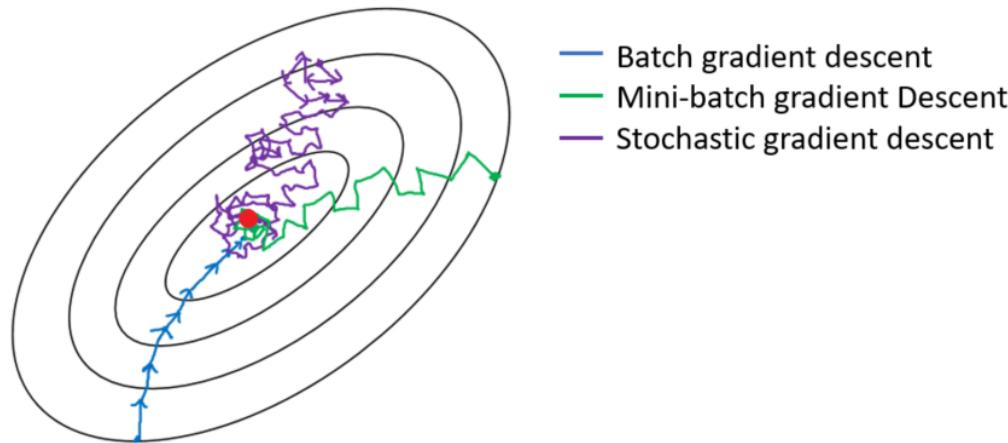
- iterative algorithm for finding a minimum of a function



How do we minimise the loss function?

Stochastic Gradient Descend (SGD)

- most used for deep learning
- uses only a small data sample to calculate the gradient
 - less precision
 - much faster to compute (remember: we do this iteratively)



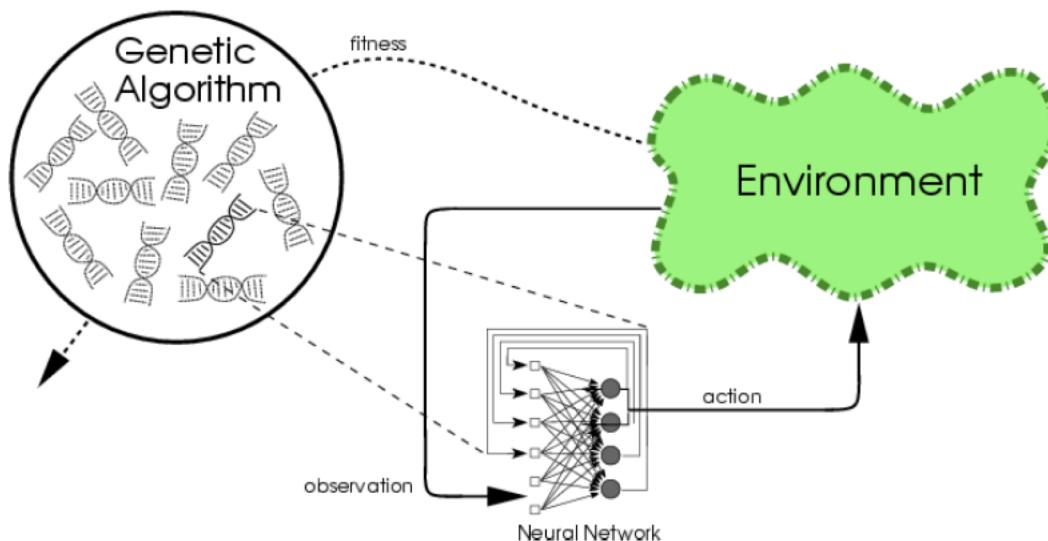
Neural Networks are...

- **AWSOME!**
- Finding the gradient involves some calculus ☹
 - kidding, that's not an excuse of a data scientist ☺



Neuroevolution

- **training NNs through evolution!**
- example of a **gradient-less training method**
- **simple genetic algorithm** can sometimes outperform modern RL algorithms
- can be faster due to good parallelization

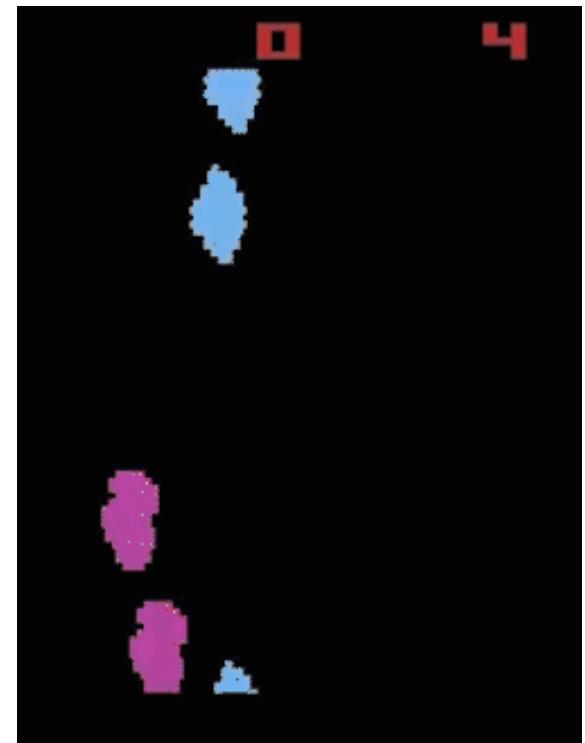


Genetic algorithms as a competitive alternative for training deep neural networks

Paper by Felipe Petroski Such, Vashisht Madhavan, Edoardo Conti, Joel Lehman Kenneth O. Stanley and Jeff Clune from Uber AI labs

<https://arxiv.org/abs/1712.06567>

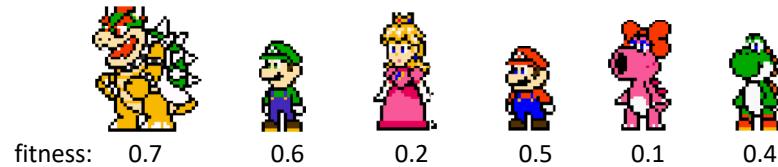
- DNNs trained to play **Atari games from pixels**
 - benchmark
- compares well to traditional RL algorithms
- **novel approach to store the genomes**
- largest neural network ever evolved with a traditional evolutionary algorithm





Principle of the Genetic Algorithm:

1. evaluate each individual to get the fitness score:



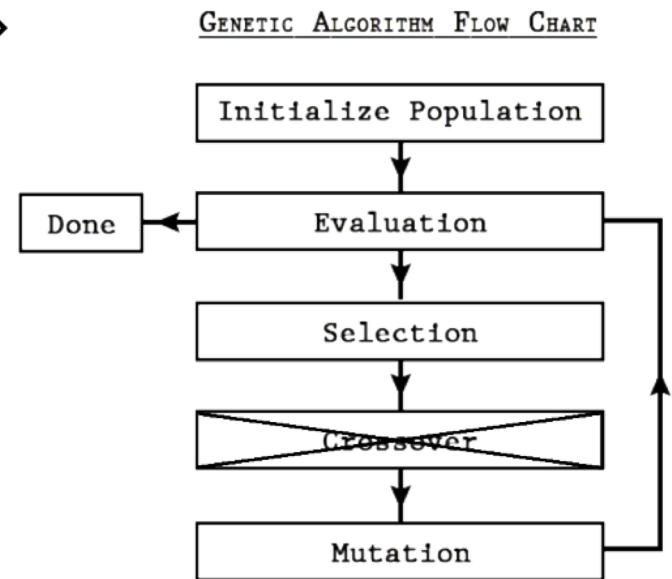
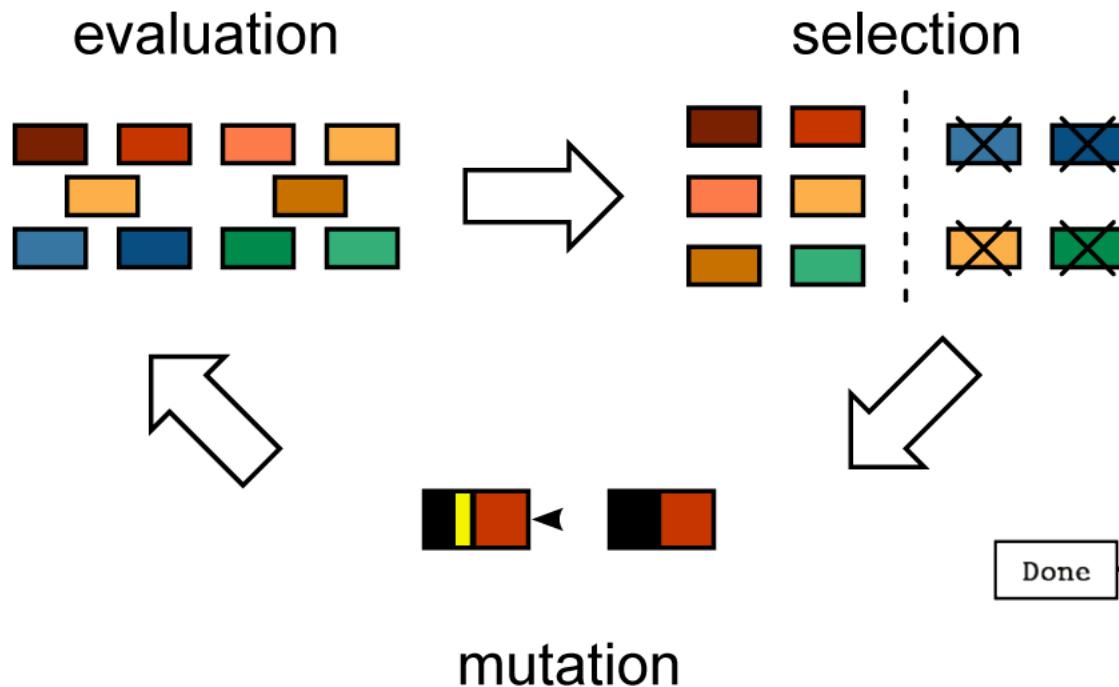
2. remove the low performing ones



3. produce offsprings with random mutations



No crossover involved!



New approach to store the



!

- traditional approach - store each individual θ as s vector of weights
 - takes up a lot of memory
 - does not scale well (especially in networks with large populations)
- **new approach** – store the initialisation seed and mutation seeds
 - seeds $s_1, s_2 \dots s_n$ are fed into a deterministic mutation function
 - critical innovation for an efficient implementation of a distributed deep GA

$$\theta = \psi(\text{initialisation seed}, \{s_1, s_2 \dots s_n\})$$

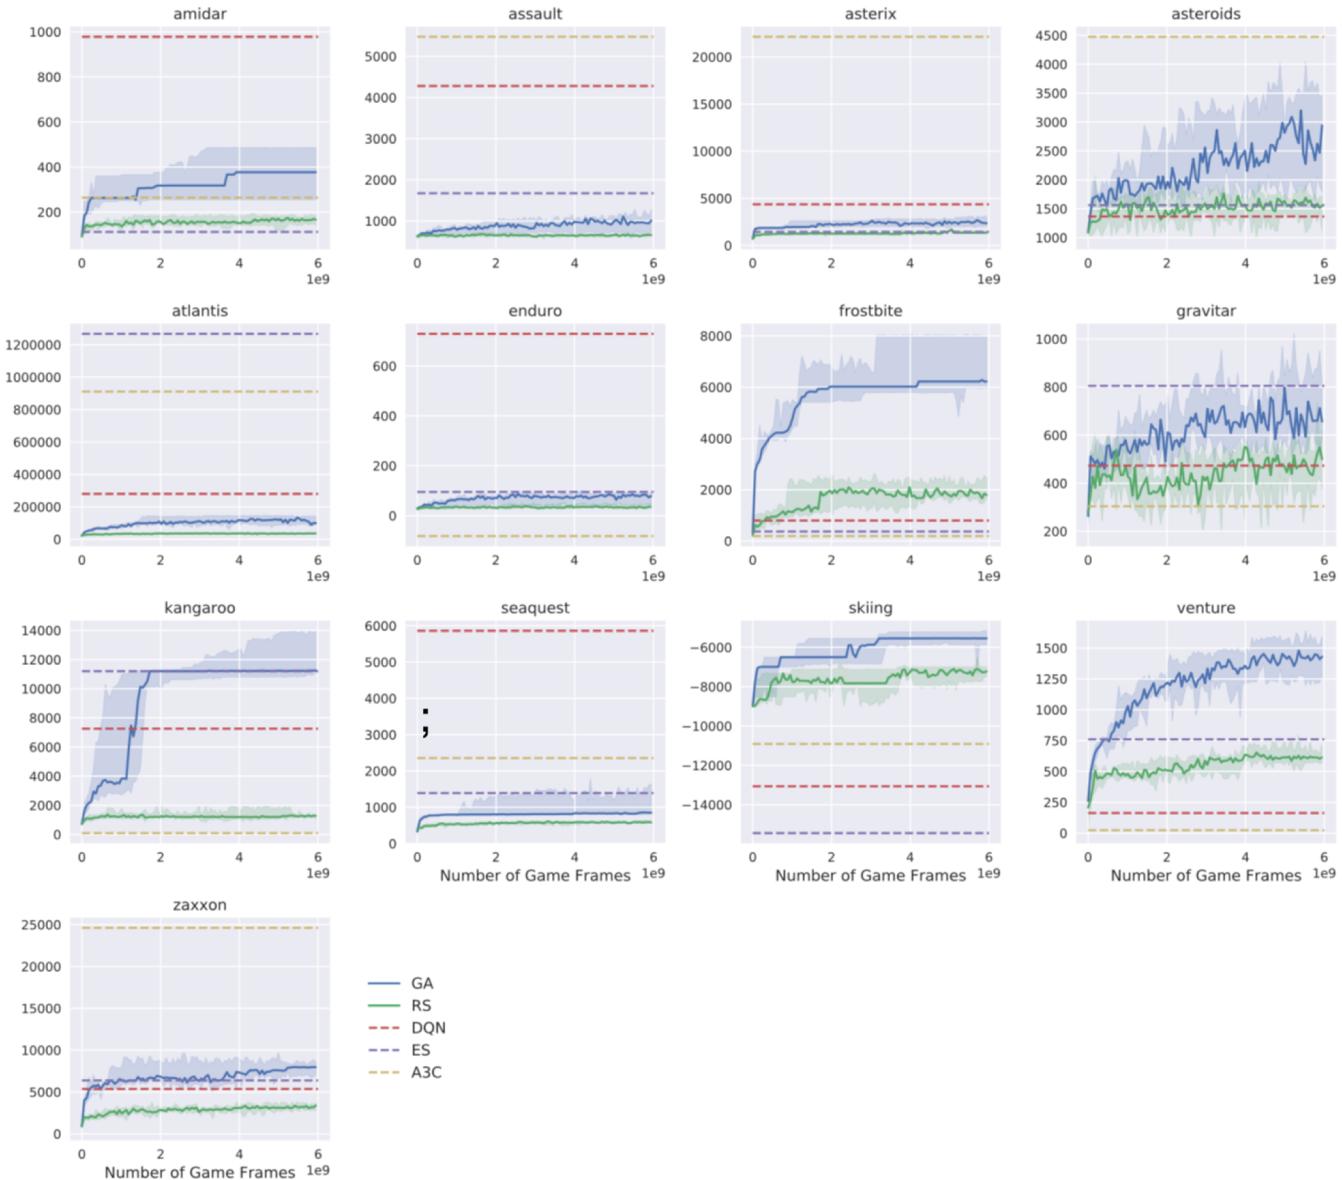
- individual θ can be easily reconstructed while being efficiently stored
 - size scales linearly with the number of generations,
 - independent from the size of network!!!

← this is awesome

Benchmarking

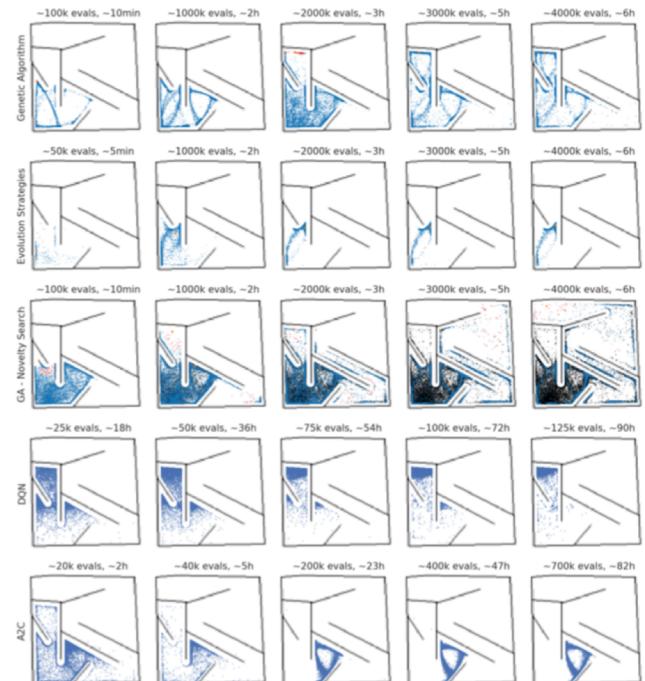
GA always outperforms random search

	DQN	ES	A3C	RS	GA	GA
Frames	200M	1B	1B	1B	1B	6B
Time	~7-10d	~ 1h	~ 4d	~ 1h or 4h	~ 1h or 4h	~ 6h or 24h
Forward Passes	450M	250M	250M	250M	250M	1.5B
Backward Passes	400M	0	250M	0	0	0
Operations	1.25B U	250M U	1B U	250M U	250M U	1.5B U
amidar	978	112	264	143	263	377
assault	4,280	1,674	5,475	649	714	814
asterix	4,359	1,440	22,140	1,197	1,850	2,255
asteroids	1,365	1,562	4,475	1,307	1,661	2,700
atlantis	279,987	1,267,410	911,091	26,371	76,273	129,167
enduro	729	95	-82	36	60	80
frostbite	797	370	191	1,164	4,536	6,220
gravitar	473	805	304	431	476	764
kangaroo	7,259	11,200	94	1,099	3,790	11,254
seaquest	5,861	1,390	2,355	503	798	850
skiing	-13,062	-15,443	-10,911	-7,679	[†] -6,502	[†] -5,541
venture	163	760	23	488	969	[†] 1,422
zaxxon	5,363	6,380	24,622	2,538	6,180	7,864



Using old algorithms developed in the neuroevolution - Novelty search

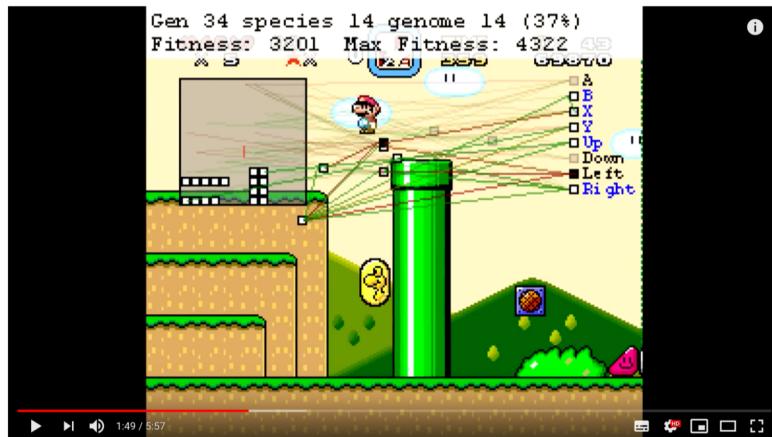
- example of using existing algorithms with DNNs
- rewarding agent for performing new behaviours – instead of fitness
 - novelty = average distance to kNN (sorted by distance) in the population/archieve
- works surprisingly well with GA
 - it can often outperform algorithms that utilize the reward signal



Conclusions

- **old algorithms** combined with **modern amounts of computing** can work surprisingly well
- neuroevolution searches differently than SGD, and thus offers an interesting alternative approach
- state-of-the-art compact **encoding technique**
 - allow for effective parallelization

Other examples of GA with NN



Marl/O - Machine Learning for Video Games

7 232 969 zhlédnutí

128 TIS. 2,3 TIS. SDÍLET ULOŽIT ...



SethBling

Publikováno 13. 6. 2015

POČET ODBĚRŮ: 1,9 MIL.

